

## 0625 Load Balancer 使用 p4 實作負載均衡器

我們現在所使用的網頁伺服器，一台機器如果性能比較好的，大概能服務一兩萬條連線，如果人數再更多，就沒有辦法負荷。所以，通常這種網頁伺服器都是用集群式的方式在服務，也就是說，它會開很多台，同時很多台在等待把使用者的請求服務分散掉，因為像一些購物節或者是搶票，使用者請求會非常多，幾台機器可能沒辦法應付，所以就有很多台機器。

### Connection hash

當很多的客戶，如果集中在一台，容易爆掉，所以負載均衡就是把客戶的請求平均分散到不同伺服器上，這樣的話每一台伺服器的壓力就比較小，反應時間快，如果全都集中在單一台就會來不及服務，在那邊等待，所以我們希望，當一鍵按下去可以快速得到我想要的東西

負載均衡器的概念就是，它會提供一個固定 ip，只要連線連進來這個固定 ip，它會自動幫你進行分發的動作

## p4-14 寫法轉換成 p4-16

執行步驟：

1. 打開終端機，切到 p4-test 資料夾
  2. 建立資料夾(mkdir hash-lb)，並切到 hash-lb 資料夾
  3. gedit hash-lb.p4 &，去 <http://csie.nqu.edu.tw/smallko/sdn/LBP4.htm> 把 load\_balance.p4 複製並貼上
  4. 在終端機輸入指令 p4c-bm2-ss -p4v 14 -pp hash-lb16.p4 hash-lb.p4
  5. gedit hash-lb16.p4 & 就把 p4-14 寫法轉換成 p4-16 了
- p4-14 轉 p4-16 網站：<https://p4tw.org/%E5%B0%87-p4-14-%E5%BF%AB%E9%80%9F%E8%BD%89%E6%8F%9B%E8%87%B3-p4-16-%E6%96%B9%E6%B3%95/>

## 負載均衡器

[Topology]

H1 (10.0.1.1) -----(P4 switch: load balancer)----- H2 (Simple HTTP Server, 10.0.2.2).  
----- H3 (Simple HTTP Server, 10.0.3.3).

Virtual IP: 10.0.0.1

1. 打開終端機，切到 `p4-test/hash-lb` 資料夾
2. `gedit commands.txt &` ，一樣去 LBP4.htm 網站把 `s1-command.txt` 複製並貼上
3. `gedit p4app.json &` ，把 `p4app.json` 貼上
4. 在終端機執行 `p4run`，開啟三個終端 `xterm h1 h2 h3`
5. 在 `h3` 執行 `python -m SimpleHTTPServer 80`
6. 在 `h2` 執行 `echo "hi" > hi.htm` 產生一個簡單的網頁，再執行 `python -m SimpleHTTPServer 80`
7. 在 `h1` 執行 `curl 10.0.0.1/hi.htm`

每次存取的時候，就會選擇不一樣伺服器

```
##### Welcome to the P4utils Mininet CLI #####  
Your P4 program is installed into the BMV2 software,  
and your initial configuration is loaded. You can interact  
with the network using the mininet CLI below.  
  
To inspect or change the switch configuration, connect to its  
CLI from your host operating system using this command:  
simple_switch_CLI --thrift-port <switch thrift port>  
  
To view a switch log, run this command from your host:  
tail -f /home/user/p4-test/hash-lb/log/<switchname>  
  
To view the switch output pcap, check the pcap files  
/home/user/p4-test/hash-lb/pcap:  
for example run: sudo tcpdump -xxx -r s1-eth1.pcap  
  
*** Starting CLI:  
mininet> xterm h1 h2 h3  
mininet>
```

## s1-commands.txt

```
table_set_default forward nop
table_set_default ecmp_group nop
table_set_default ecmp_nhop nop
table_set_default send_frame nop
table_add forward set_nhop 10.0.1.1/32 => 00:00:0a:00:01:01 1
table_add forward set_nhop 10.0.2.2/32 => 00:00:0a:00:02:02 2
table_add forward set_nhop 10.0.3.3/32 => 00:00:0a:00:03:03 3
table_add ecmp_group set_ecmp_select 10.0.0.1/32 => 0 2
table_add ecmp_nhop set_ecmp_nhop 1 => 00:00:0a:00:02:02 10.0.2.2 2
table_add ecmp_nhop set_ecmp_nhop 2 => 00:00:0a:00:03:03 10.0.3.3 3
table_add send_frame rewrite_sip 1 => 10.0.0.1
```

原來是 0 要改成 a

如果目的地是 10.0.0.1(vip 位址)，2 代表有兩個選擇，一個是 1 一個是 2。如果 hash 的結果是 1(bit)，代表要把請求丟給伺服器 2，如果 hash 值是 3 就丟到伺服器 3(第二台)

當封包回去的時候，要把原本的來源 ip 轉換成 vip

## p4app.json

```
{
  "program": "hash-lb16.p4",
  "switch": "simple_switch",
  "compiler": "p4c",
  "options": "--target bmv2 --arch v1model --std p4-16",
  "switch_cli": "simple_switch_CLI",
  "cli": true,
  "pcap_dump": true,
  "enable_log": true,
  "topo_module": {
    "file_path": "",
    "module_name": "p4utils.mininetlib.apptopo",
    "object_name": "AppTopo"
  },
  "controller_module": null,
  "topodb_module": {
    "file_path": "",
    "module_name": "p4utils.utils.topology",
    "object_name": "Topology"
  },
  "mininet_module": {
    "file_path": "",
    "module_name": "p4utils.mininetlib.p4net",
    "object_name": "P4Mininet"
  },
  "topology": {
    "assignment_strategy": "manual",
    "default_bw": 10,
    "default_delay": "1ms",
    "auto_gw_arp": true,
    "links": [["h1", "s1"], ["s1", "h2"], ["s1", "h3"]],
    "hosts": {
      "h1": {
        "ip": "10.0.1.1",
        "gw": "10.0.1.254"
      }
    }
  }
}
```

```
},
"h2": {
  "ip" : "10.0.2.2",
  "gw": "10.0.2.254"
},
"h3": {
  "ip": "10.0.3.3",
  "gw": "10.0.3.254"
}
},
"switches": {
  "s1": {
    "cli_input": "command.txt",
    "program": "hash-lb16.p4"
  }
}
}
```

## Load\_balance.p4

```
#include <core.p4>
#include <v1model.p4>
```

```
struct meta_t {
    bit<1>  do_forward;
    bit<32> ipv4_sa;
    bit<32> ipv4_da;
    bit<16> tcp_sp;
    bit<16> tcp_dp;
    bit<32> nhop_ipv4;
    bit<32> if_ipv4_addr;
    bit<48> if_mac_addr;
    bit<1>  is_ext_if;
    bit<16> tcpLength;
    bit<8>  if_index;
}
```

```
struct mymetadata_t {
    bit<14> ecmp_select;
}
```

```
header arp_t {
    bit<16> htype;
    bit<16> ptype;
    bit<8>  hlen;
    bit<8>  plen;
    bit<16> opcode;
    bit<48> hwSrcAddr;
    bit<32> protoSrcAddr;
    bit<48> hwDstAddr;
    bit<32> protoDstAddr;
}
```

```
header ethernet_t {
    bit<48> dstAddr;
    bit<48> srcAddr;
```

```
    bit<16> etherType;
}
```

```
header ipv4_t {
    bit<4>  version;
    bit<4>  ihl;
    bit<8>  diffserv;
    bit<16> totalLen;
    bit<16> identification;
    bit<3>  flags;
    bit<13> fragOffset;
    bit<8>  ttl;
    bit<8>  protocol;
    bit<16> hdrChecksum;
    bit<32> srcAddr;
    bit<32> dstAddr;
}
```

```
header tcp_t {
    bit<16> srcPort;
    bit<16> dstPort;
    bit<32> seqNo;
    bit<32> ackNo;
    bit<4>  dataOffset;
    bit<4>  res;
    bit<8>  flags;
    bit<16> window;
    bit<16> checksum;
    bit<16> urgentPtr;
}
```

```
header udp_t {
    bit<16> srcPort;
    bit<16> dstPort;
    bit<16> length_;
    bit<16> checksum;
}
```

```

struct metadata {
    @name(".meta")
    meta_t      meta;
    @name(".mymetadata")
    mymetadata_t mymetadata;
}

```

```

struct headers {
    @name(".arp")
    arp_t      arp;
    @name(".ethernet")
    ethernet_t ethernet;
    @name(".ipv4")
    ipv4_t      ipv4;
    @name(".tcp")
    tcp_t      tcp;
    @name(".udp")
    udp_t      udp;
}

```

```

parser ParserImpl(packet_in packet, out headers hdr, inout metadata meta, inout
standard_metadata_t standard_metadata) {
    @name(".parse_arp") state parse_arp {
        packet.extract(hdr.arp);
        transition accept;
    }
    @name(".parse_ethernet") state parse_ethernet {
        packet.extract(hdr.ethernet);
        transition select(hdr.ethernet.etherType) {
            16w0x800: parse_ipv4;
            16w0x806: parse_arp;
            default: accept;
        }
    }
    @name(".parse_ipv4") state parse_ipv4 {
        packet.extract(hdr.ipv4);
        meta.meta.ipv4_sa = hdr.ipv4.srcAddr;
        meta.meta.ipv4_da = hdr.ipv4.dstAddr;
    }
}

```



```

        meta.meta.tcpLength = hdr.ipv4.totalLen - 16w20;
        transition select(hdr.ipv4.protocol) {
            8w6: parse_tcp;
            8w17: parse_udp;
            default: accept;
        }
    }
    @name(".parse_tcp") state parse_tcp {
        packet.extract(hdr.tcp);
        meta.meta.tcp_sp = hdr.tcp.srcPort;
        meta.meta.tcp_dp = hdr.tcp.dstPort;
        transition accept;
    }
    @name(".parse_udp") state parse_udp {
        packet.extract(hdr.udp);
        transition accept;
    }
    @name(".start") state start {
        meta.meta.if_index = (bit<8>)standard_metadata.ingress_port;
        transition parse_ethernet;
    }
}

```

```

control egress(inout headers hdr, inout metadata meta, inout standard_metadata_t
standard_metadata) {
    @name("_drop") action _drop() {
        mark_to_drop(standard_metadata);
    }
    @name(".rewrite_sip") action rewrite_sip(bit<32> sip) {
        hdr.ipv4.srcAddr = sip;
    }
    @name(".nop") action nop() {
    }
    @name(".send_frame") table send_frame {
        actions = {
            _drop;
            rewrite_sip;
            nop;
        }
    }
}

```

```

    }
    key = {
        standard_metadata.egress_port: exact;
    }
    size = 256;
}
apply {
    send_frame.apply();
}
}

```

```

control ingress(inout headers hdr, inout metadata meta, inout standard_metadata_t
standard_metadata) {
    @name("_drop") action _drop() {
        mark_to_drop(standard_metadata);
    }
    @name(".set_ecmp_select") action set_ecmp_select(bit<8> ecmp_base, bit<8>
ecmp_count) {
        //這個地方要來做 hash
        hash(meta.mymetadata.ecmp_select, HashAlgorithm.crc16,
(bit<14>)ecmp_base, { hdr.ipv4.srcAddr, hdr.ipv4.dstAddr, hdr.ipv4.protocol,
hdr.tcp.srcPort, hdr.tcp.dstPort }, (bit<28>)ecmp_count);
        //hash 五個欄位：來源 ip，目的 ip，通訊協定，來源埠號，目的埠號
        meta.mymetadata.ecmp_select = meta.mymetadata.ecmp_select + 14w1;
    }
    @name(".nop") action nop() {
    }
    @name(".set_ecmp_nhop") action set_ecmp_nhop(bit<48> nhop_mac, bit<32>
nhop_ipv4, bit<9> port) {
        standard_metadata.egress_spec = port;
        hdr.ipv4.dstAddr = nhop_ipv4;
        hdr.ethernet.dstAddr = nhop_mac;
        hdr.ipv4.ttl = hdr.ipv4.ttl - 8w1;
    }
    @name(".set_nhop") action set_nhop(bit<48> dmac, bit<9> port) {
        standard_metadata.egress_spec = port;
        hdr.ethernet.dstAddr = dmac;
        hdr.ipv4.ttl = hdr.ipv4.ttl - 8w1;
    }
}

```

```

}
@name(".ecmp_group") table ecmp_group {
    actions = {
        _drop;
        set_ecmp_select;
        nop;
    }
    key = {
        hdr.ipv4.dstAddr: lpm;
    }
    size = 1024;
}
@name(".ecmp_nhop") table ecmp_nhop {
    actions = {
        _drop;
        set_ecmp_nhop;
        nop;
    }
    key = {
        meta.mymetadata.ecmp_select: exact;
    }
    size = 1024;
}
@name(".forward") table forward {
    actions = {
        _drop;
        set_nhop;
        nop;
    }
    key = {
        hdr.ipv4.dstAddr: lpm;
    }
    size = 1024;
}

```

```

    apply {
        forward.apply();
        ecmp_group.apply();
        ecmp_nhop.apply();
    }
}

```

Ingress 的部分有三個 table :

**forward.apply();** : table\_add forward set\_nhop 10.0.1.1/32 => 00:00:0a:00:01:01 1

**ecmp\_group.apply();** :

table\_add ecmp\_group set\_ecmp\_select 10.0.0.1/32 => 0 2

如果他選擇 vip , 就 set\_ecmp\_select;

```

control DeparserImpl(packet_out packet, in headers hdr) {
    apply {
        packet.emit(hdr.ethernet);
        packet.emit(hdr.arp);
        packet.emit(hdr.ipv4);
        packet.emit(hdr.udp);
        packet.emit(hdr.tcp);
    }
}

```

```

control verifyChecksum(inout headers hdr, inout metadata meta) {
    apply {
        verify_checksum(true, { hdr.ipv4.version, hdr.ipv4.ihl, hdr.ipv4.diffserv,
hdr.ipv4.totalLen, hdr.ipv4.identification, hdr.ipv4.flags, hdr.ipv4.fragOffset,
hdr.ipv4.ttl, hdr.ipv4.protocol, hdr.ipv4.srcAddr, hdr.ipv4.dstAddr },
hdr.ipv4.hdrChecksum, HashAlgorithm.csum16);
        verify_checksum_with_payload(true, { hdr.ipv4.srcAddr, hdr.ipv4.dstAddr,
8w0, hdr.ipv4.protocol, meta.meta.tcpLength, hdr.tcp.srcPort, hdr.tcp.dstPort,
hdr.tcp.seqNo, hdr.tcp.ackNo, hdr.tcp.dataOffset, hdr.tcp.res, hdr.tcp.flags,
hdr.tcp.window, hdr.tcp.urgentPtr }, hdr.tcp.checksum, HashAlgorithm.csum16);
    }
}

```

```

control computeChecksum(inout headers hdr, inout metadata meta) {
    apply {
        update_checksum(true, { hdr.ipv4.version, hdr.ipv4.ihl, hdr.ipv4.diffserv,
hdr.ipv4.totalLen, hdr.ipv4.identification, hdr.ipv4.flags, hdr.ipv4.fragOffset,
hdr.ipv4.ttl, hdr.ipv4.protocol, hdr.ipv4.srcAddr, hdr.ipv4.dstAddr },
hdr.ipv4.hdrChecksum, HashAlgorithm.csum16);
        update_checksum_with_payload(true, { hdr.ipv4.srcAddr, hdr.ipv4.dstAddr,
8w0, hdr.ipv4.protocol, meta.meta.tcpLength, hdr.tcp.srcPort, hdr.tcp.dstPort,

```

```
hdr.tcp.seqNo, hdr.tcp.ackNo, hdr.tcp.dataOffset, hdr.tcp.res, hdr.tcp.flags,  
hdr.tcp.window, hdr.tcp.urgentPtr }, hdr.tcp.checksum, HashAlgorithm.csum16);  
    }  
}
```

```
V1Switch(ParserImpl(), verifyChecksum(), ingress(), egress(), computeChecksum(),  
DeparserImpl()) main;
```